



Proof-of-concept · status update

The problem, in one sentence

We constantly **rewrite** deep-learning code to make it run faster.

But how do we *know* the faster version computes the **same thing** as the original?

- A rewrite that is *almost* identical can silently change the results.
- Silently wrong math → wrong predictions, bugs that are very hard to find.

We want a way to be **sure**, not hopeful.

Our example: the “QKV” step inside every Transformer

Attention models — the engine behind modern AI — build three things, **Q**, **K**, **V**, from the input X using three weight matrices.

The clear way (unfused) — three separate matrix multiplications:


$$Q = X W_q \quad K = X W_k \quad V = X W_v$$

The fast way (fused) — glue the weights together, do **one** big multiply, then split:

$$W = [W_q \mid W_k \mid W_v] \quad QKV = X W \quad \rightarrow \text{split into } Q, K, V$$

These **should** give identical results. The question: do they — *always*?

The two programs, side by side



ESBMC proves these two produce **the same result for every input we model**.

We run the program on a **handful of example inputs** and compare.

Like tasting a few spoonfuls of soup and hoping the whole pot is fine.

- Testing only covers the inputs you happened to try.
- A bug hiding in some other input is simply **missed**.
- **Passing tests $\neq$ correct.**

What we actually want: a proof, not a sample

Instead of trying a few inputs, check **every input we model** — **all at once** — mathematically.

Two possible outcomes:

- ✓ **PROVEN equal** — guaranteed identical for *every* input in the model.
- ✗ **Counterexample** — one specific input where they differ, handed straight to you.

Either way, you learn something **certain**.

ESBMC is an automated *model checker*.

- 1 It **reads the program** and turns the question “*are these two always equal?*” into one giant logic puzzle.
- 2 An automated **solver** then either finds an input that breaks it, or proves that **no such input exists** (within the model).

Fully automated once the program is encoded — you write *no* proofs by hand.

The trick that turns a test into a proof

Ordinary test	Our verification
one random tensor <code>torch.randn(...)</code>	a symbolic tensor = <i>any</i> numbers (bounded)
checks that one case	checks the whole modelled range at once
“looks fine”	a guarantee — within the model

We assert `unfused == fused`, feed it *any* input in range, and let ESBMC settle it.

Inputs are bounded to a finite range to exclude “not-a-number”. (→ Assumptions, backup.)

Result: the example is PROVEN correct

- ✓ Proven equal for **all inputs we model** — and **bit-for-bit exact** (IEEE-754), not a “close enough” tolerance.
- ✓ Done **two ways**: a plain-math encoding *and* the real `torch.mm + torch.allclose`.
- ✓ When we **deliberately break it** (swap two columns), ESBMC **catches it** and shows the exact failing input.

What “modelled” means (fixed sizes, bounded inputs) → **Assumptions & Scope**, backup.

What the suite covers (and what it doesn't — yet)

6 / 6 checks behave as expected — small *on purpose*: it is breadth, not scale.

Dimension	Coverage
Problem classes	QKV projection · bias-fused linear ($XW + b$)
Encodings	plain-math and torch-native (<code>torch.mm</code> / <code>allclose</code>)
Check type	a proof (clean $\Rightarrow$ ✓) and a refutation (bug $\Rightarrow$ ✗ + counterexample)

Every case is double-checked: we prove the correct version *and* confirm a broken one is caught. **Early-stage PoC** — small fixed sizes for now (see Performance, backup).

- **All modelled inputs**, not a few → no hidden corner cases within the model.
- **Exact** → catches tiny numerical drift that testing would wave through.
- **Automated** → no hand-written proofs.
- **Actionable failures** → it gives you the precise input that breaks things.

Concretely: **detect incorrect optimizations early** and **validate fusion / compiler transformations** *before* they ship.

Real PyTorch-style code pushed ESBMC into territory it could not yet handle — so we **added what was missing** (4 contributions merged upstream):

- A torch **operational model** — `mm`, `matmul`, `cat`, `split`, `allclose`.
- Fixes so **matrix math over nested lists works** — values/types survive function returns and copies.
- A **floating-point element-type fix** for matrices built by comprehension (we found, root-caused, and fixed it).

Result: PyTorch-style code is now verifiable in ESBMC **at all** — reusable by anyone.

- ✓ Motivating **QKV** example **fully supported** (two encodings, exact proof).
- ✓ Extended to a **second** class: bias-fused linear ($XW + b$).
- → **Bigger matrices** and fully native fuse/split — gated on **one** performance improvement upstream (see Performance, backup).

Takeaway

We can **prove** — automatically, for **every modelled input** — that a fast PyTorch rewrite computes **exactly** the same result as the original.

Testing samples. This proves — and when it can't, it hands you the bug.

Use it to gate fusions and compiler rewrites before they ship.

Backup slides

Assumptions · Under the hood · Performance

What we prove

- Equivalence of the two programs *as written*, for **all inputs in the modelled range**.
- Floating-point modelled **exactly** (IEEE-754, bit-precise) — not real arithmetic, not a hand-waved tolerance.

What we model / bound

- **Fixed tensor sizes** (bounded model checking; concrete, e.g. $S=1, D=2, H=1$).
- **Inputs bounded** to a finite range, which excludes NaN/Inf.

What we do not claim: arbitrary/symbolic sizes · performance equivalence · behaviour outside the bounds.

- 1 ESBMC **unrolls** the program (fixed sizes $\Rightarrow$ finite) and turns each assert into a **verification condition**.
- 2 Tensors are modelled as lists of **IEEE-754 floats**; `torch.mm` / `torch.allclose` have operational models.
- 3 The condition is essentially: *for all bounded inputs*, `fused[i][j] == unfused[i][j]`.
- 4 An **SMT solver** (Bitwuzla / Z3) decides it: **UNSAT** $\Rightarrow$ **proved**; **SAT** $\Rightarrow$ **counterexample**.

```
X, Wq, Wk, Wv = symbolic, bounded           # any numbers in range
assert torch.allclose(unfused, fused)         # -> one SMT formula over FP
```

Check	Time (Bitwuzla, macOS arm64)
scalar QKV proof	~11 s
torch-native QKV proof	~110 s
refuting a broken variant	31–292 s

- Cost grows quickly with tensor size (more multiply-adds → a bigger FP formula).
- `torch.cat` / `torch.split` are modelled but **unwinding-heavy** (a $1 \times 1 \times 3$ `cat+split` ≈ 2.5 min) → we fuse/split by manual indexing for now.
- **Scaling to realistic sizes is the main open bottleneck** (tracked upstream).